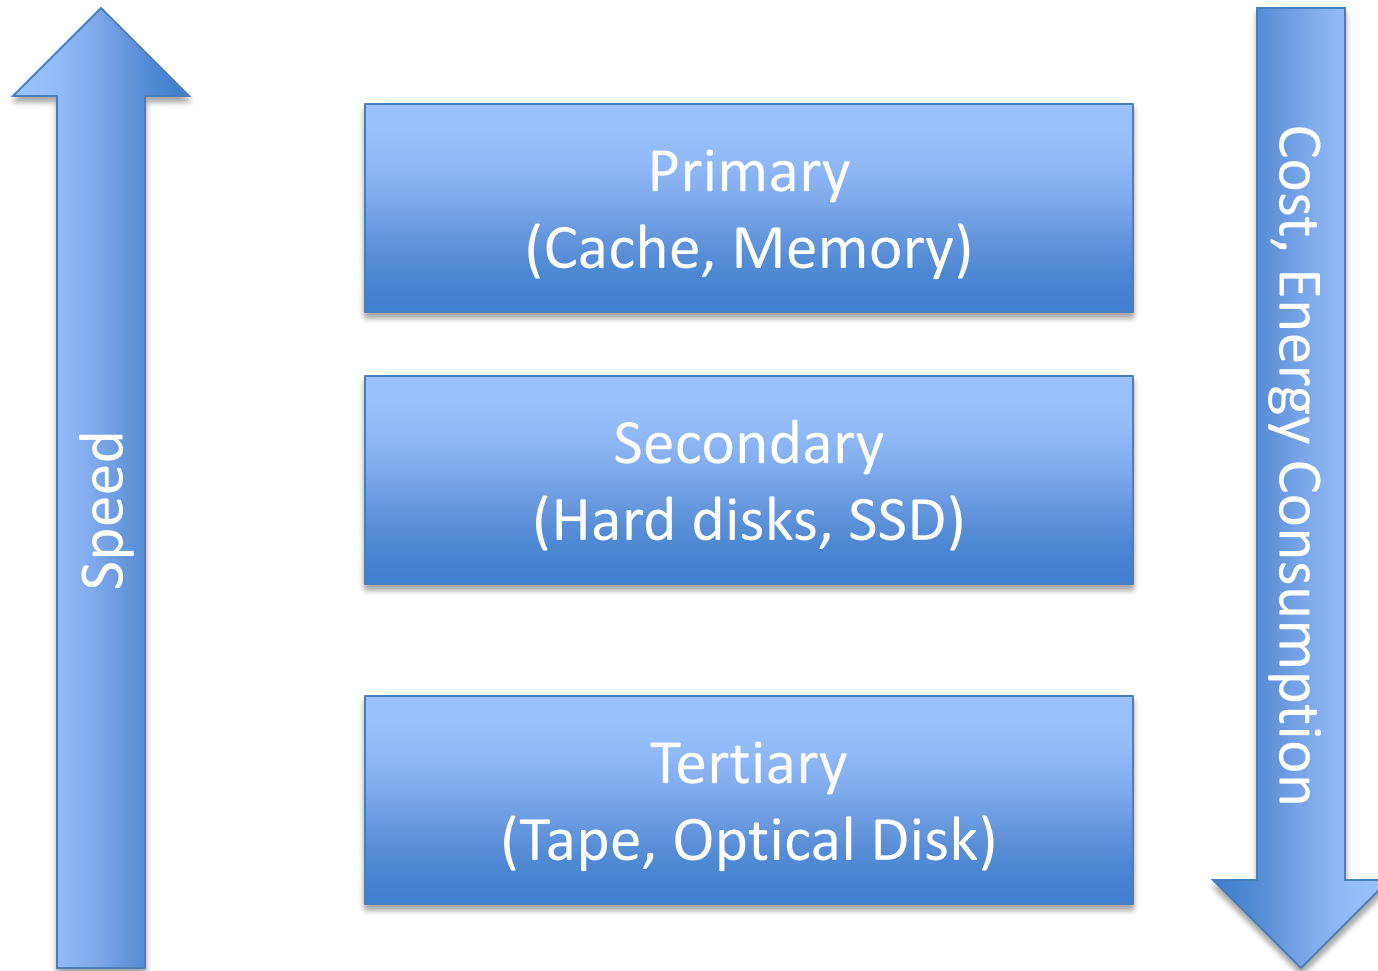


Storage Systems

INF 551

Wensheng Wu

Storage hierarchy



Storage operations: CRUD

- We should think about storage in terms of storage operations
- CRUD:
 - (C)reate
 - (R)ead
 - (U)pdate
 - (D)elele

Storage is not forever

- Transient vs. permanent Errors
- Media failures
 - One part of the storage media stops working
 - Cannot read from/write to physical location in media
- Device failures
 - Failure of entire device, data may not be recoverable
- Detecting failure
 - Periodic scanning of media
 - On-device monitoring
 - Inability to access device

Characterizing a storage device

- Capacity (bytes)
 - How much data it can hold
- Cost (\$\$\$)
 - Price per byte of storage
- Bandwidth (bytes/sec)
 - Number of bytes that can be transferred per second
 - Note that read and write bandwidth may be different
- Latency (secs)
 - Time between initiating a request and an action
 - In the case of storage, action is to deliver 1st Byte

Time to complete an operation

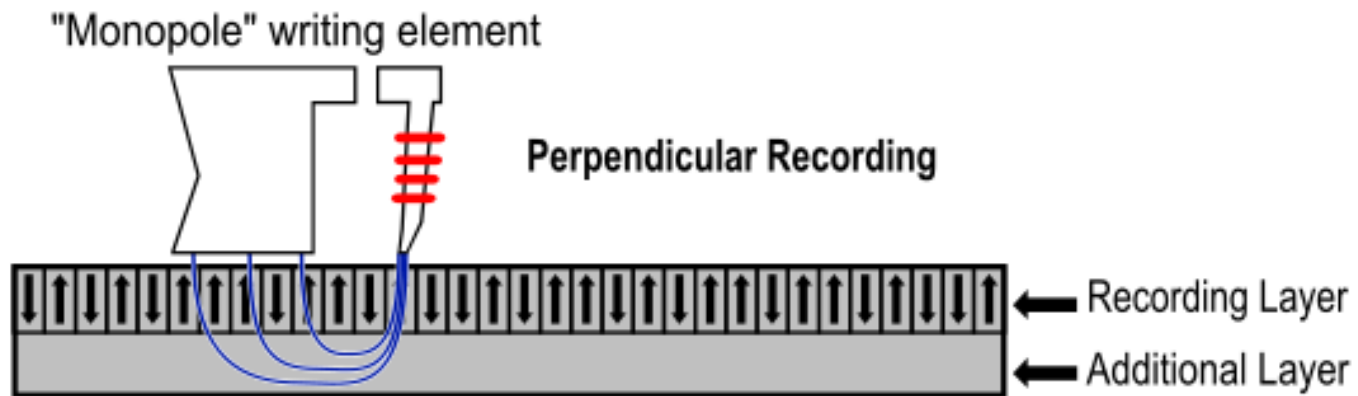
- Time to complete an operation depends on both bandwidth and latency
 - $\text{CompletionTime} = \text{Latency} + \text{Size}/\text{Bandwidth}$
- The time of a work load will depend on
 - Technology, e.g., hard drive/ssd
 - Operation type, e.g., read/write
 - Number of operations in the work load
 - Access pattern (random vs. sequential)

Access pattern

- Sequential
 - Data to be accessed are located next to each other on the device
- Random
 - Access data located randomly on storage device


Magnetic recording

- Write head
 - Applies current to write head
 - Changes direction of magnetic field under head



Reading

- Read head senses direction of magnetic field



Road map

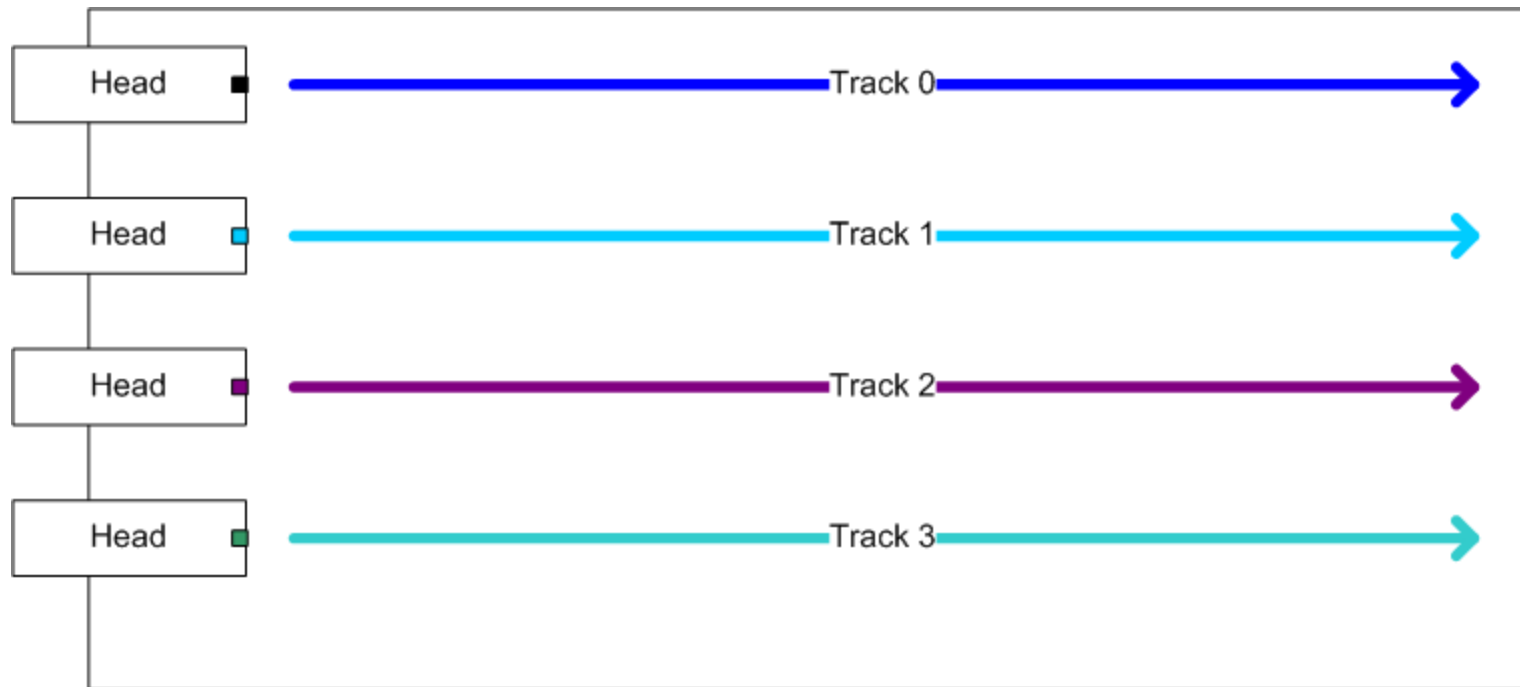
- Tapes



- Hard disk
- Disk scheduling algorithms
- Solid state drive

Linear tape

- Data recorded on parallel tracks that span the length of the tape



Tapes

- Current technology is LTO
 - Linear Tape-Open (an open standard)
- Characteristics
 - Capacity up to 6.25 TB per tape (LTO-7)
 - Drive cost ~ \$2500
 - Tape cost ~ \$45 for 2.5TB tape
- Tape access time (~ minute)
 - Time to mount the tape
 - Time to wind the tape to correct position
- Data transmission rates ~ 250MB/sec



↖
\$200
300MB/s
LTO-7

Performance characteristics

- High latency/low cost makes tape most appropriate for “archival” storage
 - Low frequency of reads
 - Very large data objects
- Random access will be slow due to latency
 - Sequential reads will be fast

Linear tape file system


- Two partitions on tape
 - First contains metadata and directories. Tape reader can find and load this very quickly
 - Second contains blocks for data
- Directory structure coded in XML
 - Self describing file format...



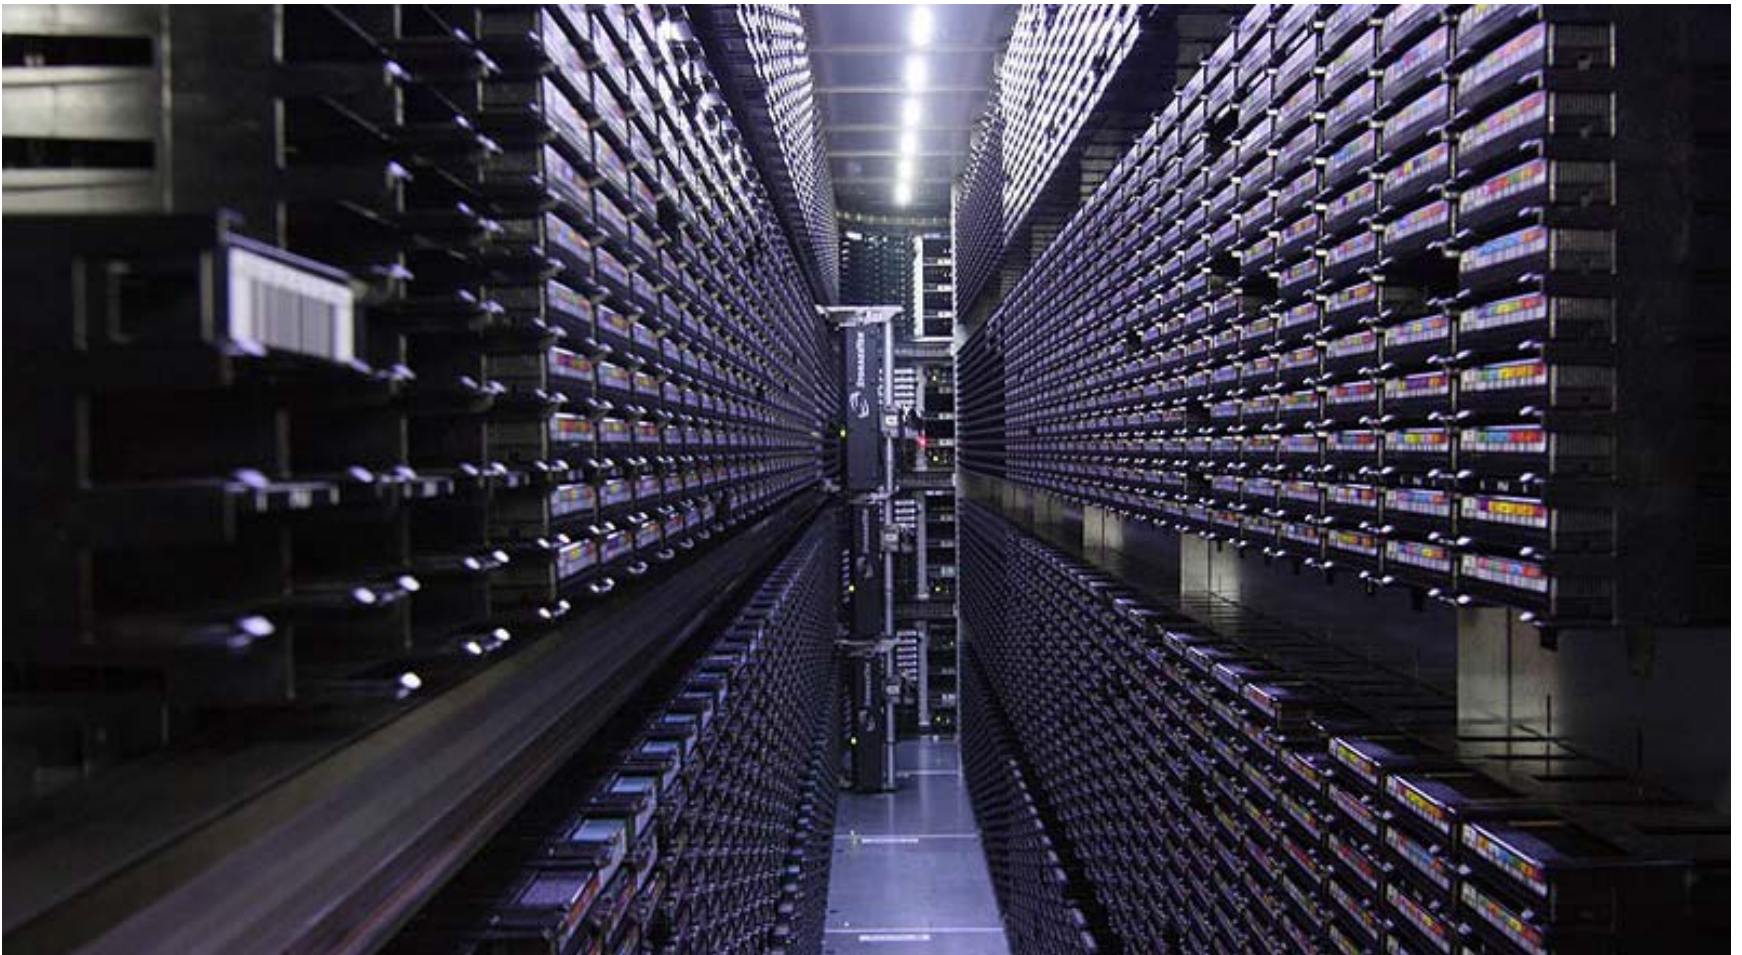
Tape Cartridge



Raw Capacity in GB



A tape library




Inside a robotic tape library

- <https://www.youtube.com/watch?v=nYfTtvpQ778>



Road map

- Tapes
- Hard disk 
- Disk scheduling algorithms
- Solid state drive

Hard disk drives

- Perhaps the most pervasive form of storage
- Basic Idea:
 - One or more spinning magnetic platters
 - Typically two surfaces per platter
 - Disk arm positions over the radial position (tracks) where data is stored
 - It swings across tracks (but do not extend/shrink)
 - Data is read/written by a read/write head as platter spins



Internal of hard disk

Actuator

Spindle




Platter

Disk head

Disk arm and platter




Disk head close-ups



Disk head close-ups





Disk head movement

- Hard disk head movement while copying files between two folders (e.g., partition c to d)
 - <https://www.youtube.com/watch?v=BlB49F6ExkQ>



2GB Storage in 1980s (\$250,000)




Physical characteristics

- 3.5" (diameter, common in desktops)
- 2.5" (common in laptops)
- Rotational speed
 - 5,400 RPM
 - 7,200 RPM
 - 4,800 RPM
 - 10,000 RPM (6ms/rotation)
- Between 5-7 platters
- Current capacity up to 10TB (Western Digital)

Disk organization


- Disk has multiple tracks
- Each track is divided into N fixed size sectors
 - Typical sector size is 512 bytes (old) or 4KB (new)
 - Sectors can be numbered from 0 to $N-1$
 - Entire sector is written “atomically”
 - All or nothing

A simple disk drive (one track only)



Rotational latency

- Waiting for the right sector to rotate under the head
 - On average: $\frac{1}{2}$ of time for a full rotation
 - Worst case?
 - Best case?




Rotation time

- Assume 10,000 RPM (rotations per minute)

$$\frac{\text{Time (ms)}}{1 \text{ rotation}} = \frac{60000 \text{ ms}}{1000 \text{ rotation}} = \frac{6 \text{ ms}}{\text{rotation}}$$

Multiple tracks: add seek times



Average seek time is about $1/3$ max seek time

Transmission time

- Assume that we transfer 512KB
- Assume 128 MB/sec transmission bandwidth
 - $512\text{KB}/128\text{MB} * 1000\text{ms}$
 $= 4\text{ms}$

Completion time

- $T = T_{\text{seek}} + T_{\text{rotation}} + T_{\text{transfer}}$
 - T_{seek} : Time to get the disk head on right track
 - T_{rotation} : Time to wait for the right sector to rotate under the head
 - T_{transfer} : Time to actually transfer the data

Example

- Capacity 4TB
- # platters: 4
- # heads: 8
- Bytes per sector: 4096
- Transmission bandwidth: 100MB/sec
- Maximum seek time: 12ms
- RPM: 10,000

Time to transfer a file

- The file occupies 100 sectors (sequentially)
- Avg. seek time =?
- Avg. rotational latency =?
- Transfer time = ?

Sector vs. block

- Block = 1 or more sectors
- Disk typically transfers one block at a time
 - Why?
- We will assume one block = one sector
 - Unless stated otherwise

Sequential operations

- Assume all sectors involved are on same track
 - We may need to seek to the right track
 - And rotate to the first sector
- But no rotation/seeking is needed afterward

Sequential vs. random

- Consider disk with 7ms avg seek, 10,000 RPM platter speed and 50 MB/sec transfer rate, 4KB/block
- Sequential access of 10 MB
 - Completion time = $7 + 3 + 10/50 * 1000 = 210\text{ms}$
 - Actual bandwidth = $10\text{MB}/210\text{ms} = 47.62 \text{ MB/s}$
- Random access of 10 MB (2,500 blocks)
 - Completion time = $2500 * (7 + 3 + 4/50) = 25.2\text{s}$
 - Actual bandwidth = $10\text{MB} / 25.2\text{s} = .397 \text{ MB/s}$

Road map

- Tapes
- Hard disk
- Disk scheduling algorithms
- Solid state drive



Scheduling Problem

- Multiple requests in a queue
- Schedule the requests
 - So that they are served as quickly as possible
 - E.g., measured by avg. # of tracks/cylinders travelled per request
- Handled by either OS or disk controller or both

Example


- Queue: 120, 30, 60, 135, 80, 20, 95, 160, 70, 5
- Head on 50
- Innermost track: 0
- Outermost track: 199

Scheduling Algorithms

- FCFS: first-come first-served
- SSTF: shortest seek time first
- SCAN (Elevator), C-SCAN (circular SCAN)
- LOOK, C-LOOK (circular LOOK)

FCFS (first-come first-served)

- Queue: 120, 30, 60, 135, 80, 20, 95, 160, 70, 5




FCFS

- # of tracks travelled: 675 (avg 67.5/request)
 - 50 -> 120 = 70
 - 120 -> 30 = 90
 - ...
- Problems:
 - A lot of head movements
 - Going back and forth

SSTF (shortest seek time first)

- # of tracks travelled: $160 - 50 + 160 - 5 = 265$
- Is this solution optimal?




SSTF problems

- Starvation possible unless queue is frozen
- E.g., 160 was just served, but more requests near 160, say 150 or 170, arrive
 - Requests 5, 20, 30 may need to wait for long time

SCAN (Elevator)

- Seek to closest track (assume idle now)
 - continue to the end
 - then reverse the direction
- # of tracks travelled: $199 - 50 + 199 - 5 = 343$



SCAN problems

- Travel right back to the area just served
- Middle area has more chances to be served
 - Worst case for request on edge: ~ 400 tracks
 - Worst case for middle area: ~200 tracks

SCAN problems


- Starvation still possible if queue is not frozen
 - E.g., right after serving 120, more requests on the **same** track come

C-SCAN

- Scan in only one direction:
 - from innermost (track 0) to outermost track (199)
- Jump back to innermost when reaching end
 - 0 → 199
 - Jump to 0
 - 0 → 199
 -

Example

- # of tracks travelled : $199 - 50 + 199 + 30 = 378$




Observations

- More uniform treatments on all tracks
 - Worst case for request on edge: 200 tracks + jump time
- But time is wasted on jump
 - although jumping is faster than moving one track at a time
 - Only one acceleration and one deceleration


LOOK

- Similar to SCAN, but scan to last request only
- # of tracks travelled : 265
 - Behave the same as SSTF in this example



C-LOOK

- Similar to C-SCAN, but scan to last request
 - Jump to the last request (the opposite side) too
- # tracks travelled: $160 - 50 + 160 - 5 + 30 - 5 = 290$




Comparisons

- FCFS: 675
- SSTF: 265
- SCAN: 343
- C-SCAN: 378
- LOOK: 265
- C-LOOK: 290

Comparisons

- Lightly-loaded system
 - SSTF and LOOK are good choices
- Heavily-loaded system
 - SCAN and C-SCAN are better
 - SSTF more likely now to suffer from starvation
 - Saving from LOOK likely less significant now

Road map

- Tapes
- Hard disk
- Disk scheduling algorithms
- Solid state drive 

Solid State Drive



Chips



Solid State Drives

- All electronic, made from flash memory
- Lower energy consumption than hard drive
- Significantly more expensive, less capacity
 - About a factor of 10 more expensive
- Limited lifetime, can only write a limited number of times.
 - E.g., 100, 000 write cycles for SLC (single-level cell) memory

Solid State Drives

- Same form-factor and control interface as magnetic disks
- Significantly better latency
 - No seek or rotational delay
- Consistent bandwidth for sequential & random:
 - Benefits from improved latency
 - However, writes take significantly longer than reads


Writing to SSD is complicated

- Can not overwrite a page
 - Need to erase its block (at certain point) instead
- SSD controllers take care of all these details

SSD

- Contains a number of flash memory chips
 - Chip -> dies -> planes -> blocks -> pages (rows) -> cells
 - Cells are made of floating-gate transistors
- Page is the smallest unit of data transfer between SSD and main memory
 - Much like a block in hard disk

Die Layout



Dies, planes, block, and pages

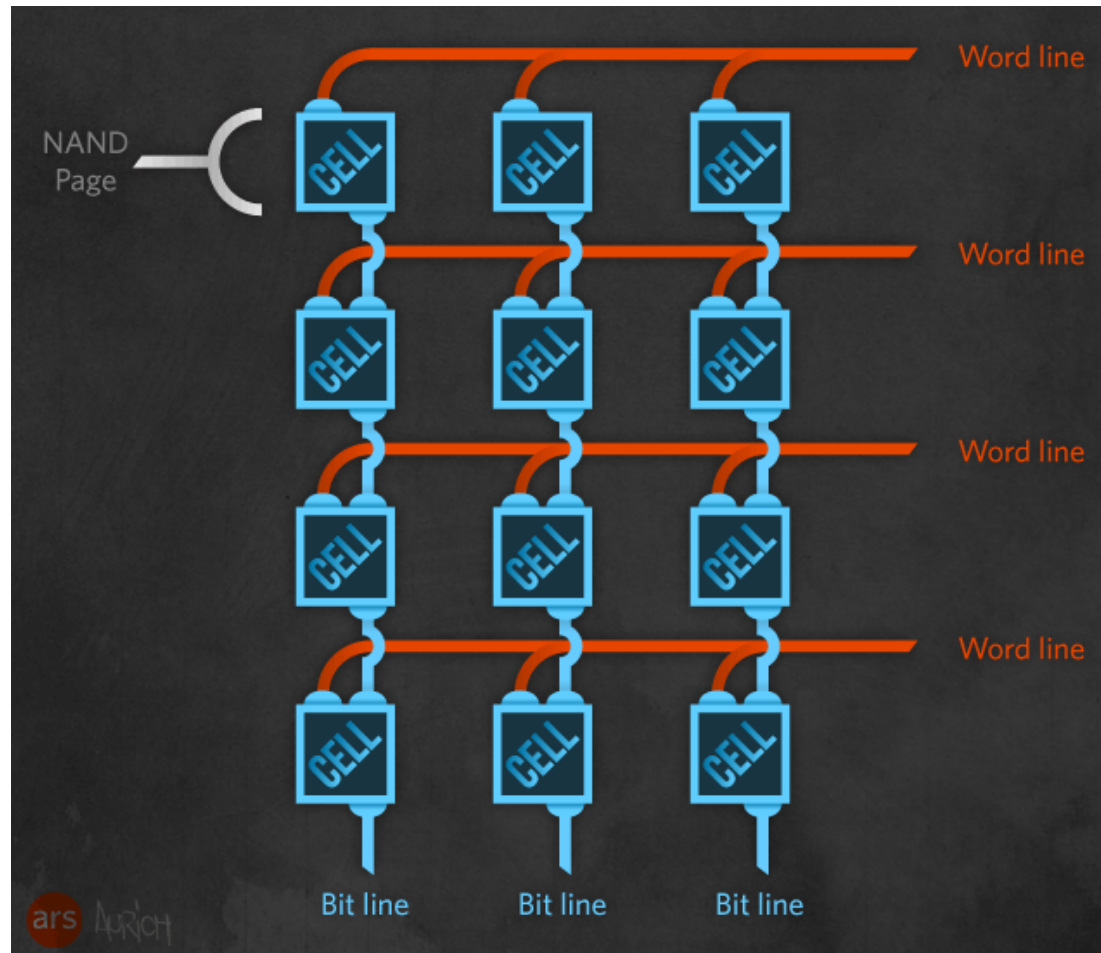
- Typically, a chip may have 1, 2, or 4 dies
- A die may have 1 or 2 planes
- A plane has a number of blocks
 - Block is the smallest unit that can be erased
- A block has a number of pages
 - Page is the smallest unit that can be programmed/written to

Typical page and block sizes

- Common page sizes: 2K, 4K, 8K, and 16K
- A block typically has 128 to 256 pages


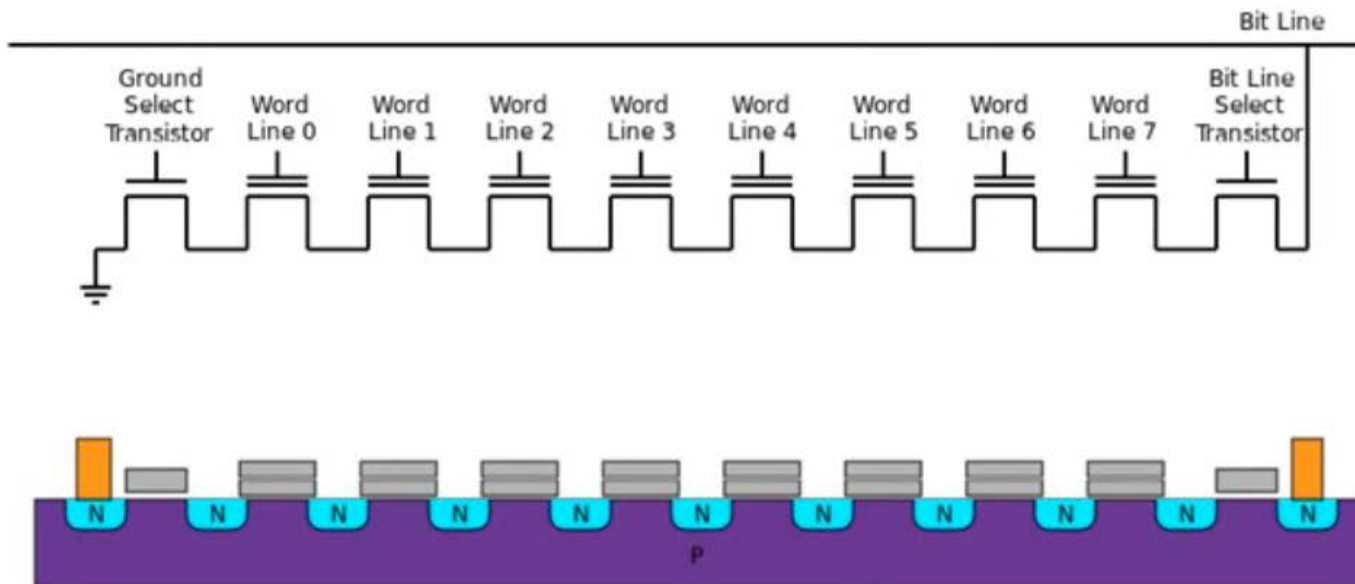
=> Block size: 256KB to 4MB

NAND flash

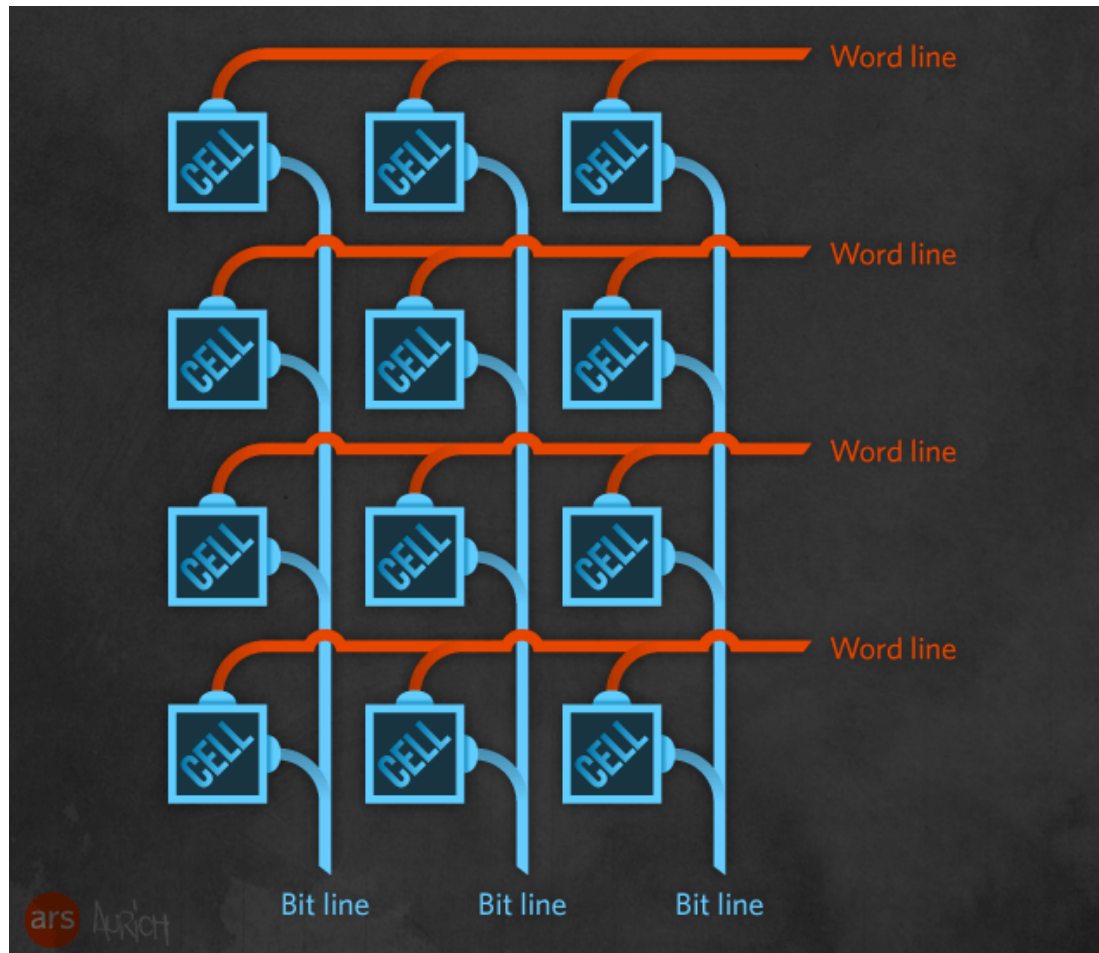


NAND flash layout

- Transistors are strung together in a series
 - Similar to the transistors in an NAND gate





NOR flash layout



NOR flash layout

- Floating gate transistors are wired in parallel
 - Each is directly connected to bit line (also ground)
 - Similar to transistors in a NOR gate (to output F)




NAND vs NOR

- Bit line
 - NOR has individual bit line (for cell), so circuit more complex
 - NAND ties up all bit lines, save space, to allow larger capacity
- Default value of cell
 - NOR: 0
 - **NAND: 1**


Write vs. erase

- **Page** is the smallest unit that can be read or **written** (also called programmed)
- **Block** is the smallest unit that can be **erased**
 - i.e., make cells "empty" (storing default values)



Operation	Area
Read	Page
Program (Write)	Page
Erase	Block

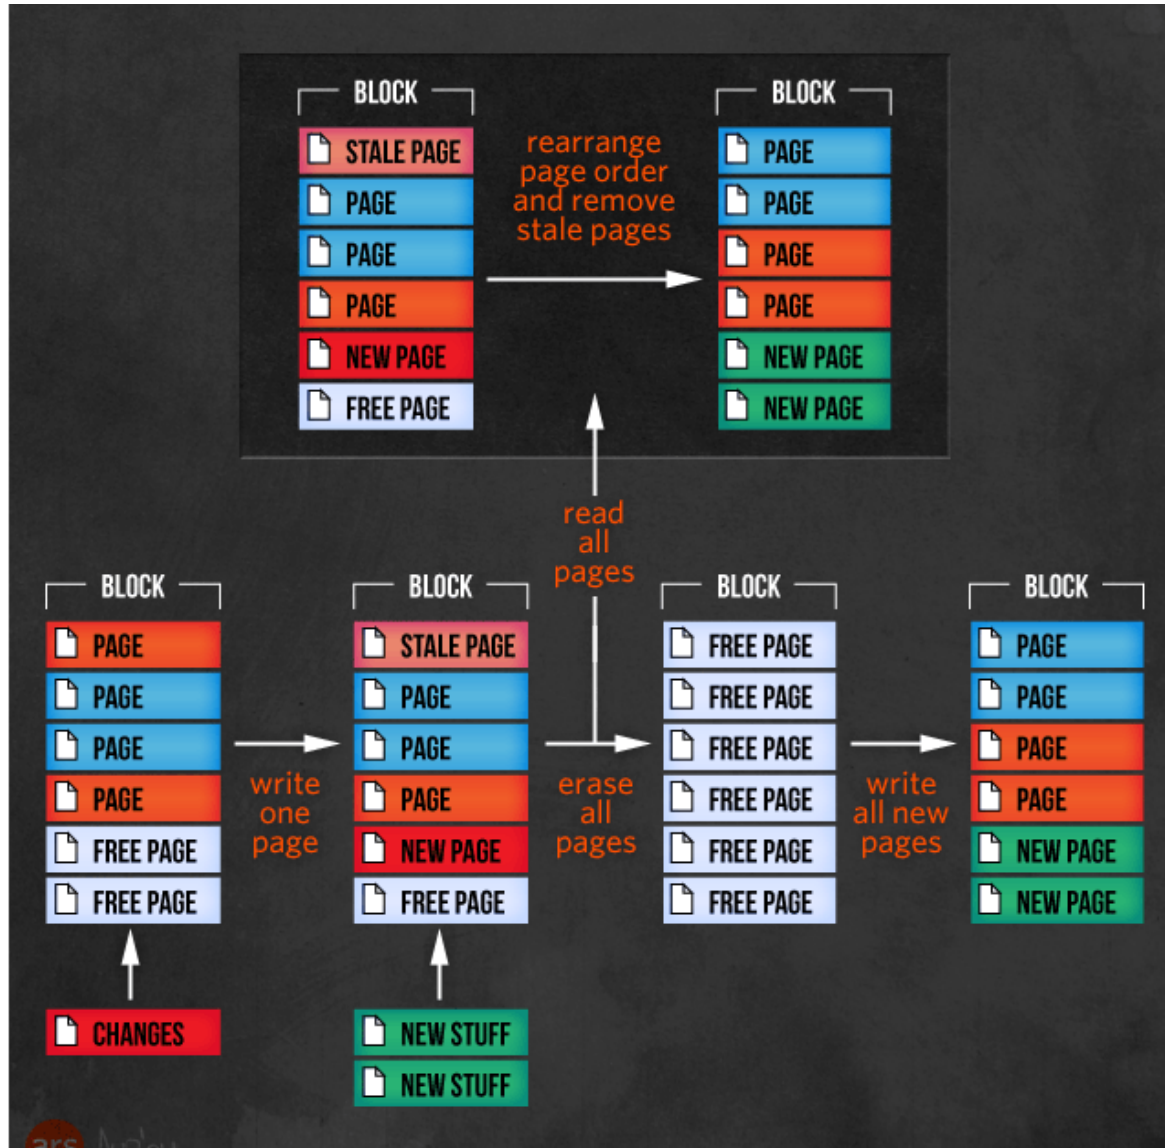
Floating gate transistor



Write vs. overwrite (NAND flash)

- Write: $1 \Rightarrow 0$
 - Need to apply high voltage to the gate
- Overwrite: $0 \Rightarrow 1$
 - Need to apply much higher voltage than write
 - May stress surrounding cells
 - So dangerous to do on individual pages

Example



Latencies: read, write, and erase

	SLC	MLC	TLC	HDD	RAM
P/E cycles	100k	10k	5k	*	*
Bits per cell	1	2	3	*	*
Seek latency (μ s)	*	*	*	9000	*
Read latency (μ s)	25	50	100	2000-7000	0.04-0.1
Write latency (μ s)	250	900	1500	2000-7000	0.04-0.1
Erase latency (μ s)	1500	3000	5000	*	*
<i>Notes</i>	* metric is not applicable for that type of memory				
<i>Sources</i>	P/E cycles [20] SLC/MLC latencies [1] TLC latencies [23] Hard disk drive latencies [18, 19, 25] RAM latencies [30, 52] L1 and L2 cache latencies [52]				

P/E cycle

- P: program/write
- E: erase
 - Every erase damages oxide layer surrounding the floating-gate to some extent
- P/E cycle:
 - Data are written to cells (P): cell value from 1 -> 0
 - Then erased (E): 0 -> 1

Read more

- [Solid-state revolution: in-depth on how SSDs really work](#)
- [How do SSDs work?](#)
 - <http://www.extremetech.com/extreme/210492-extremetech-explains-how-do-ssds-work>

References

- How Flash Memory Works
 - <https://www.youtube.com/watch?v=msi5GDz9JIw>
- Floating Gate Basics
 - <http://www.cse.scu.edu/~tschwarz/coen180/LN/flash.html>
- Friend of Flash
 - http://www.nnc3.com/mags/LM10/Magazine/Archive/2008/86/040-041_logfs/article.html