



INF 551

Wensheng Wu

Installation on EC2

- Create a new yum repository file for MongoDB
 - `cd /etc/yum.repos.d`
 - `sudo vi mongodb-org-3.4.repo`
- Add the following lines to the file:
 - `[mongodb-org-3.4]`
 - `name=MongoDB Repository`
 - `baseurl=https://repo.mongodb.org/yum/amazon/2013.03/mongodb-org/3.4/x86_64/`
 - `gpgcheck=1`
 - `enabled=1`
 - `gpgkey=https://www.mongodb.org/static/pgp/server-3.4.asc`

Installation on EC2

- `sudo yum -y install mongodb-org`
- `sudo service mongod start`
 - Start the server
- `sudo service mongod stop`
 - Stop it

Installation on EC2

- For more details see instructions here:
 - <https://docs.mongodb.com/v3.4/tutorial/install-mongodb-on-amazon/>

Document store

- MongoDB is a document database
- A document is similar to a JSON object
 - Consists of field-value pairs
 - Value may be another document, array, string, number, etc.
- Document = record/row in RDBMS

Collections

- Documents are stored in a collection
- Collection = table in RDBMS
- But documents may have different structures
 - In contrast, records in RDBMS have the same schema

Primary key

- Every document has a unique `_id` field
 - That acts as a primary key

MongoDB shell

- mongo

```
[ec2-user@ip-172-31-52-194 ~]$ mongo
MongoDB shell version: 3.2.10
connecting to: test
Server has startup warnings:
2016-11-10T22:46:56.897+0000 I CONTROL [initandlisten]
2016-11-10T22:46:56.897+0000 I CONTROL [initandlisten] ** WARNING: /sys
/kernel/mm/transparent_hugepage/defrag is 'always'.
2016-11-10T22:46:56.897+0000 I CONTROL [initandlisten] **          We sug
gest setting it to 'never'
2016-11-10T22:46:56.897+0000 I CONTROL [initandlisten]
> |
```


Create a new database

- No need to explicitly create it, just use it
 - It will be automatically created once you add a collection (i.e., table) to it

```
> show databases;
local 0.000GB
> use inf551
switched to db inf551
> show databases;
local 0.000GB
> use inf551
switched to db inf551
> db.createCollection('person')
{ "ok" : 1 }
> show databases;
inf551 0.000GB
local 0.000GB
```

```
> use inf551
switched to db inf551
> show collections
person
> show tables
person
> |
```

Databases

- use inf551
 - Switch to database "inf551"
- show databases
 - List all databases
- show tables
 - List all tables/collections in the current db
 - Can also say "show collections"

Create/drop a collection

- `db.createCollection('person')`
 - `db` is a shell variable representing the current db
- `db.person.drop()`
 - Dropping a collection

Adding documents

- `db.person.insert({"_id": 1, "name": "john smith"})`
- `db.person.insert({"_id": 1, "name": "david smith"})`
 - Error: duplicate key!

ObjectId()

- ObjectId() function creates an ID
- `db.person.insert({"_id": ObjectId(), "name": "john smith"})`

```
writeResult({ "nInserted" : 1 })
> db.person.find()
{ "_id" : 1, "name" : "john smith" }
{ "_id" : ObjectId("58250aec7c61126eba98db48"), "name" : "john smith" }
```

ObjectId()

- `db.person.insert({"name": "john smith"})`
 - Here no specification of `"_id"` field
 - But an id will be automatically created

```
> db.person.find()
{ "_id" : 1, "name" : "john smith" }
{ "_id" : ObjectId("58250aec7c61126eba98db48"), "name" : "john smith" }
{ "_id" : ObjectId("58250d56249e740a9ddfbacc"), "name" : "john smith" }
> |
```

ObjectId()

- A 12-byte hexademical value
 - E.g., 58250aec7c61126eba98db48
- Among 12 bytes:
 - 4-byte: the seconds since 1970/1/1
 - 3-byte: machine identifier
 - 2-byte: process id
 - 3-byte: a counter, starting with a random value

Embedded sub-document

- `db.person.insert(
 {
 "name": "david johnson",
 "address": {"street": "123 maple",
 "city": "LA",
 "zip": 91989},
 "phone": ["323-123-0000", "626-124-0999"]
 })`

Array

Insert some more documents

- `db.person.insert({"name": "kevin small", "age": 35})`
- `db.person.insert({"name": "mary lou", "age": 25})`

Query

- `db.person.find()`
 - Return all documents in person
- `db.person.find({"name": "kevin small"})`
 - Return all documents with specified name


Using query operators

- `db.person.find({"age": {$gt: 25}})`
- `db.person.find({"name": "kevin small", "age": {$gt: 25}})`
 - Specify "and" condition
- `db.person.find({ $or: [{"name": "kevin small"}, {"age": {$gt: 25}}] })`
 - Specify "or" condition

Query operators

- Introduced by \$
- \$lt, \$gt, \$lte, \$gte, \$ne
 - Comparison operators
- \$or, \$and, \$not
 - Logical operators

Projection

- `db.person.find(
 {"age": {$ne: 25} },
 {"name":1, "age": 1}
)`

1: included in result; 0: do not
- This will return name and age (plus `_id`)
 - i.e., similar to 'select `_id`, name, age from users where age != 25'

Projection

- `db.person.find(
 {"age": {$ne: 25} },
 {"name":1, "age": 1, "_id": 0}
)`
- This does not return id, e.g.,
 `{ "name" : "john smith" }`
 `{ "name" : "david johnson" }`
 `{ "name" : "kevin small", "age" : 35 }`

Example

- Without projection

```
> db.person.find({"age": 25})
{ "_id" : ObjectId("582559b19f185cd8ccf23ff6"), "name" : "mary lou", "age" : 25 }
```

- With projection

```
> db.person.find({"age": 25}, {"name": 1, _id: 0})
{ "name" : "mary lou" }
```

Update documents

- `db.person.update(`
 `{ "age": { $gt: 25 } },`
 `{ $set: { "status": "C" } },`
 `{ multi: true }`
)
- Existing documents may not have status field; if not, insert it instead
- Update one or all documents

Similar to:

Update users set status = 'C' where age > 25

Another example

- `db.person.update({}, {$set: {"status":'C'}}, {multi:true})`
 - Note the empty query `{}`
- Add "status" field to all documents

Remove fields

- `db.person.update({}, {$unset: {"status": ""}}, {multi: true})`
- Remove the "status" field from all documents

Delete

- `db.person.remove({})`
 - Remove all documents
- `db.person.remove({ "age": {$gt: 30} })`
 - Delete with a condition

Count()

- `db.person.count()`
 - Return # of documents in the person collection

Query on embedded document

- Using **dot notation** to identify field in embedded document
- `db.person.find({"address.city": "LA"})`
 - Return all documents whose city sub-field of address field = "LA"

Aggregation

- `db.person.aggregate([{"$group": {_id: "$age", total:{$sum:1}}}])`

`{ "_id" : 25, "total" : 1 }`

`{ "_id" : 35, "total" : 1 }`

`{ "_id" : null, "total" : 4 }`

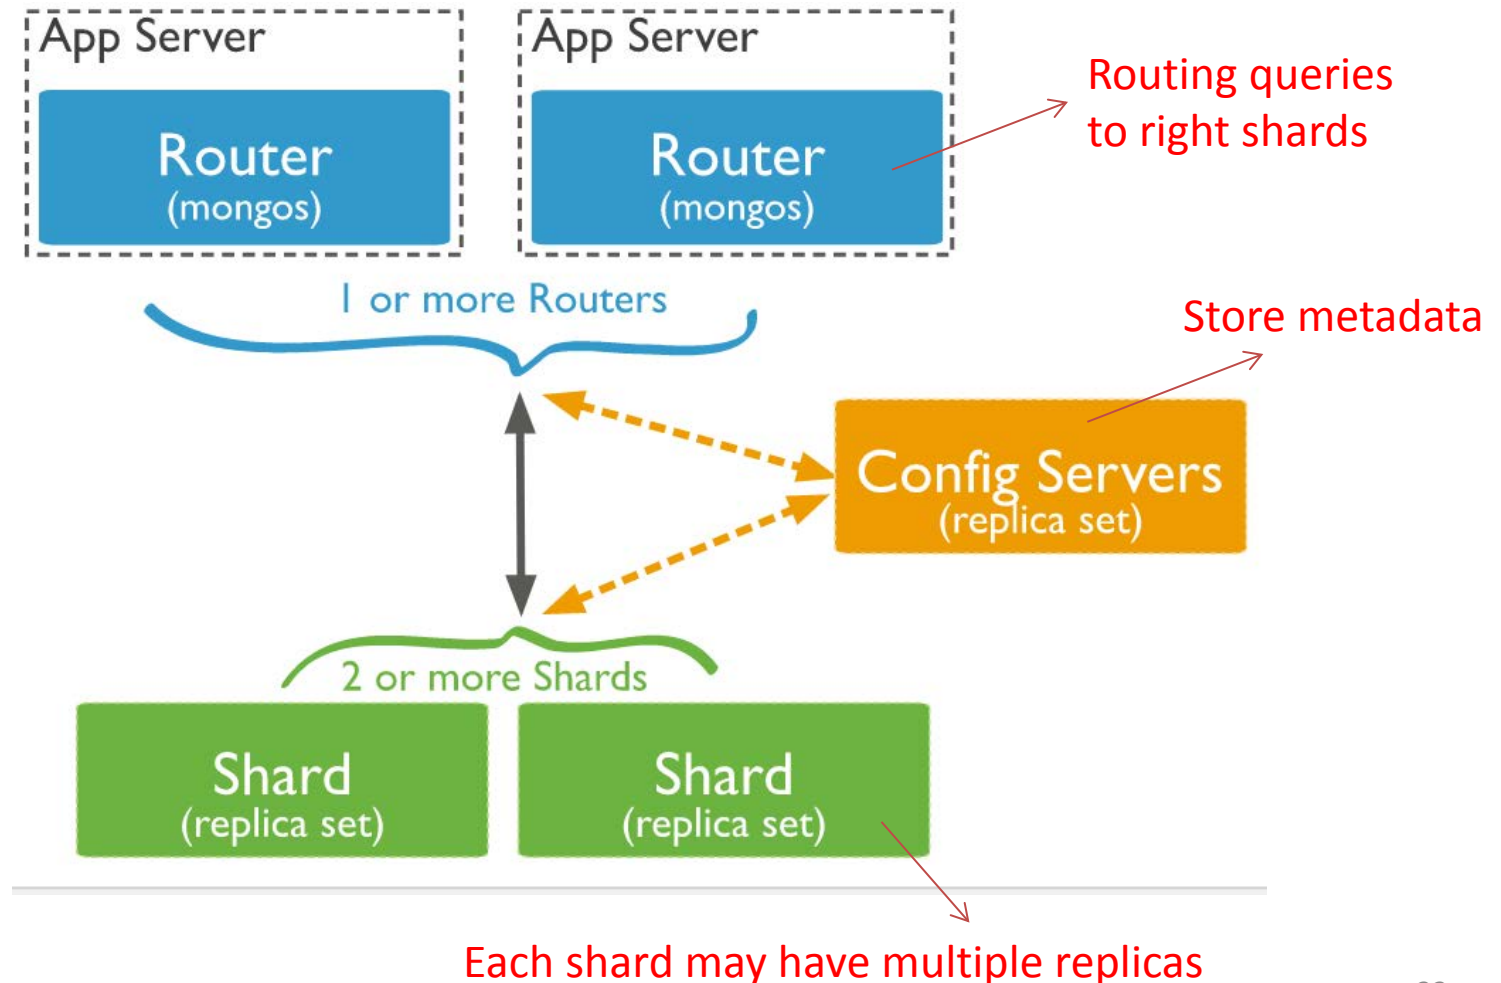
Sharding in MongoDB

- Done at collection-level
 - Distribute records in a collection over multiple machines
- User can specify shard key
 - i.e., a field in a document
- Support sharding by key range or hashing

Sharding


- Method of distributing data across multiple machines
- Sharding for horizontal scaling

Shared cluster



Sharded and non-sharded collections

- Collection 1 is sharded; collection 2 is not



Import external dataset

- `mongoimport --db inf551 --collection lax --file lax.json`
 - No need to pre-create inf551 and lax if they do not exist yet
- More details:
 - <https://docs.mongodb.com/getting-started/shell/import-data/>

Resources

- CRUB operations in MongoDB
 - <https://docs.mongodb.com/v3.4/crud/>
 - Create => insert()
 - Read => find()
 - Update => update()
 - Delete => remove()
- Cursor in mongo shell
 - <https://docs.mongodb.com/v3.3/tutorial/iterate-a-cursor/>
 - <https://docs.mongodb.com/v3.4/reference/method/js-cursor/>